

HACKATHON — AI/WORKS™ 2026

Confidence map

Visible reasoning for confident delivery

The invisible risk in software delivery

Engineering teams don't fail for lack of tools — they fail because uncertainty stays hidden



Ambiguous specifications

Teams start building from specs that contain hidden gaps, contradictions, and assumptions — without knowing it.



Blind architectural decisions

SLAs are committed, infrastructure is chosen, timelines are set — without validating feasibility against real constraints.



Late discovery is expensive

A risk found in production is a sprint lost, an incident, or a component that excludes users — none of it visible at design time.


Current AI tools — Copilot, ChatGPT, Cursor — **generate answers** but do not make uncertainty visible. They don't show what they're assuming. They produce output; they don't expose the reasoning.

THE GAP

No tool today tells a team **what it doesn't know** before the first line of code is written.

Six specialized agents. One system.

Each agent reasons independently — a consolidator acts as a neutral LLM-as-a-Judge across all outputs



The tool in action

NovaBank · International payments — real output from the analysis engine

- **Green:** Explicitly defined in the specification — safe to implement
- **Yellow:** Reasonably inferred — validate before building
- **Red:** High uncertainty — do not implement without clarification

22%

Global confidence score
The spec needs at least
one more refinement pass

△ TOP CRITICAL FINDING

"SLA contradiction: 2 s target vs synchronous legacy gateway with undefined latency"

CONFIDENCE MAP — NODE VIEW



Three hackathon challenges. One platform.

Confidence Map was designed to address all three areas directly



Spec-to-Data Pipelines

Unstructured PRDs and user stories are transformed into structured, queryable findings with explicit confidence levels, categories, and recommended actions — ready for JIRA or any downstream tool.

Structured · Traceable · Actionable



LLM-as-a-Judge

Six specialized agents act as independent evaluators of different domains. A consolidator agent then acts as a neutral second-opinion judge across all outputs — eliminating single-model bias.

6 agents · 1 consolidator · Bias-free



Harness engineering

Phased orchestration with concurrency control, per-finding guardrails that validate output quality, TTL-based result store, correlation IDs for full traceability, and real-time SSE streaming — all running on Claude.

Retry · Guardrails · Observability

"It's not QA of the code. It's QA of the idea."

Visible reasoning for confident delivery

Evans Ladera & Daniel Santibanez · Thoughtworks

Python 3.12 · FastAPI

Next.js 15 · React Flow

Claude Sonnet 4.6

6 AI agents

94 tests · 90% coverage